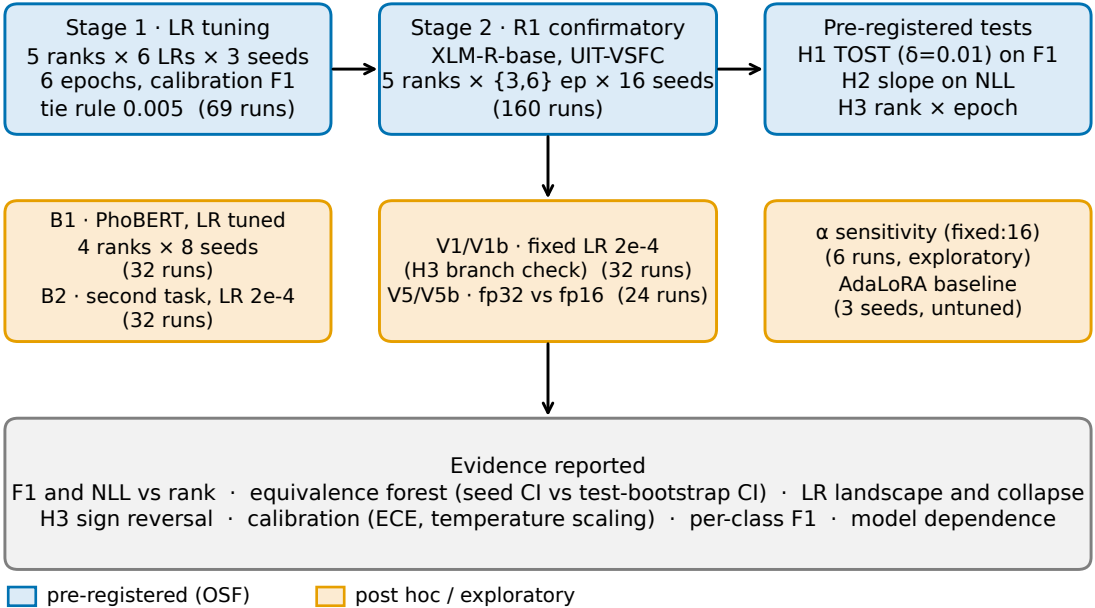
Stage 1 · LR tuning
5 ranks × 6 LRs × 3 seeds
6 epochs, calibration F1
tie rule 0.005 (69 runs)
Stage 2 · R1 confirmatory
XLM-R-base, UIT-VSFC
5 ranks × {3,6} ep × 16 seeds
(160 runs)
Pre-registered tests
H1 TOST (δ=0.01) on F1
H2 slope on NLL
H3 rank × epoch
B1 · PhoBERT, LR tuned
4 ranks × 8 seeds
(32 runs)
B2 · second task, LR 2e-4
(32 runs)
V1/V1b · fixed LR 2e-4
(H3 branch check) (32 runs)
V5/V5b · fp32 vs fp16 (24 runs)
α sensitivity (fixed:16)
(6 runs, exploratory)
AdaLoRA baseline
(3 seeds, untuned)
Evidence reported
F1 and NLL vs rank · equivalence forest (seed CI vs test-bootstrap CI) · LR landscape and collapse
H3 sign reversal · calibration (ECE, temperature scaling) · per-class F1 · model dependence
pre-registered (OSF)
post hoc / exploratory